

RentCompass

A UK Student Rental Recommendation Agent

A Complete Architectural Specification

This document explains every design detail of the agent: how it thinks, how it remembers, how a message is routed to the right capability, and how its answers are kept truthful. It assumes a reader who is comfortable with mathematics and probability but new to language-model agents and to software engineering. Each mechanism is first stated intuitively, then specified exactly.

A LangGraph state-machine agent with layered memory,
deterministic routing, and a grounding critic

Contents

1	What this system is, in one page	3
2	Mental model: state machines, LLMs, and tools	3
2.1	A language model is a function	3
2.2	A tool is also a function — but a truthful one	4
2.3	A state machine ties them together	4
3	The state: the object that flows through the graph	4
3.1	Reducers: how updates are merged	5
3.2	Field-by-field specification	5
4	The graph: nodes, edges, and the shape of a turn	7
4.1	The nine core nodes	7
4.2	Static vs. dynamic edges	8
4.3	Three optional nodes	8
5	Routing: how a message reaches the right capability	9
5.1	The priority ladder	9
5.2	The LLM classifier	10
5.3	Reference resolution	10
6	The bounded agent loop and web-research fan-out	10
6.1	Loopable tools and the reflect controller	11
6.2	Three hard stops — all failing closed	11
6.3	The map–reduce fan-out for web research	11
7	Tools: the agent’s hands	12
7.1	The uniform contract	12
7.2	The registry	12
7.3	Idempotency: making retries safe	13
7.4	search_properties in depth	13
7.5	The model router	14
8	Memory: what the agent remembers, and for how long	15
8.1	The three time-scales	15
8.2	The layered long-term memory	15
8.3	Memory as tools	16
9	From HTTP request to agent turn: the web and session layer	16
9.1	The request lifecycle	16
9.2	Composing the query: layered context injection	17
9.3	Two isolation axes, two persistence mechanisms	17
9.4	Concurrency: snapshot–compute–commit	18
10	Ranking listings: a multi-objective utility	18
10.1	Hard constraints, then soft scoring	18
11	The grounding critic: catching hallucinated numbers	19
11.1	What “grounded” means, precisely	19
11.2	Correction, not deletion	19

12 Safety, security, and honesty guards	19
13 A worked example: one turn end-to-end	20
Glossary	21

1 What this system is, in one page

RentCompass is a conversational assistant that helps international students find rental housing in the United Kingdom. A user types a message — in English or Chinese — such as “*find me a studio near UCL under £1500*” or “这个房源附近安全吗” (“is the area around this listing safe?”). The system must decide what the user *wants*, gather *real* data (live rental listings, crime statistics, transit fares, maps), reason over it, and reply in the user’s own language — while never inventing a listing that does not exist and never leaking its own instructions.

The core idea. A large language model (LLM) is very good at understanding language and writing fluent prose, but it is unreliable when asked to *decide what to do* and it will happily *fabricate* facts. So we do *not* let the LLM run free. Instead we wrap it in a **state machine**: a fixed graph of small steps (“nodes”), where deterministic code (ordinary if-statements and regular expressions) makes the high-stakes control decisions, and the LLM is called only for the two things it is genuinely good at — *classifying intent* and *writing the final answer from data we hand it*. Every LLM output is then checked by more deterministic code before the user ever sees it.

This architecture is built on **LangGraph**, a library for expressing an agent as a directed graph over a shared *state* object. The rest of this document is organised as follows:

- Section 2 builds the mental model: state machines, LLMs, and tools.
- Section 3 defines the *state* — the single data structure that flows through the graph — and its update rules (“reducers”).
- Section 4 walks the *graph*: every node, every edge, and the exact flow of a turn.
- Section 5 is the heart of the system: how a message is *routed* to a capability, mixing deterministic guards with an LLM classifier.
- Section 6 explains the *bounded agent loop* (the “reflect” controller) and the map–reduce fan-out for web research.
- Section 7 catalogues the *tools* — the agent’s hands — and how they are registered, versioned, and made idempotent.
- Section 8 is the *memory* system: short-term vs. long-term, per-conversation vs. per-user, and the Generative-Agents scoring that decides what to recall.
- Section 9 covers the *web/session layer*: how an HTTP request becomes a turn, and how many users and many conversations stay isolated.
- Section 10 gives the *multi-objective ranker* that scores listings, and Section 11 the *grounding critic* that catches hallucinations.
- Section 12 collects the *safety and security* guards.

2 Mental model: state machines, LLMs, and tools

2.1 A language model is a function

For our purposes an LLM is a function

$$f_{\theta} : (\text{text prompt}) \mapsto (\text{text completion}),$$

parameterised by fixed weights θ , that returns a plausible continuation of its input. It has two properties we must engineer around:

1. **Stochasticity.** Sampling is controlled by a *temperature* $T \geq 0$. At $T = 0$ the output is (nearly) deterministic — the most likely token each step; at $T > 0$ it is diverse. We choose T per task: $T = 0$ for classification and grounded generation, $T > 0$ only where diversity helps.
2. **No ground truth.** f_θ has no access to the current rental market, today’s crime data, or live transit fares. Anything it “knows” is frozen in θ and may be wrong or invented. Therefore *facts must come from tools*, and the LLM’s job is to *synthesise*, not to recall.

The concrete models are served through DeepSeek’s OpenAI-compatible API and chosen per role by a small *router* (Section 7.5).

2.2 A tool is also a function — but a truthful one

A *tool* is a piece of ordinary code the agent may call, e.g. “look up the crime score for postcode SW8” or “scrape live listings near UCL”. Formally a tool is a function

$$g : (\text{typed parameters}) \mapsto \text{ToolResult},$$

where `ToolResult` is a uniform envelope carrying *success/data/error* (Section 7). Tools talk to real external services — `police.uk` for crime, `OpenStreetMap` for amenities, `Transport for London (TfL)` for fares, `OnTheMarket` for listings — so their outputs are *ground truth* that the LLM must not contradict.

2.3 A state machine ties them together

The agent is a **directed graph** $G = (V, E)$. Each vertex $v \in V$ is a *node*: a function that reads the current state s and returns an update Δs . Edges determine which node runs next. Some edges are *static* (always the same successor); others are dynamic, chosen at runtime by the node itself via a `Command(goto=...)` return value. The state s is a single dictionary that every node reads and writes; `LangGraph` merges each Δs into s using per-field *reducers* (Section 3). A “turn” is one traversal of the graph from `gnodestart` to `gnodeend` for one user message.

Why not just let the LLM decide everything?

An LLM asked “which of these 12 tools should I call, with which arguments?” is right most of the time and catastrophically wrong the rest — it invents tool names, loops forever, or re-runs a fresh property search when the user only asked “is the second one safe?”. By making the *graph topology* and the *high-stakes routing decisions* deterministic, we bound the agent’s behaviour: it can only ever do things the graph permits, at most a fixed number of times, and every fact it states is checked against tool output.

3 The state: the object that flows through the graph

Everything the agent knows *during a single turn* lives in one typed dictionary, `AgentState` (defined in `src/uk_rent_agent/agent/state.py`). A node never calls another node directly; it only reads `AgentState` and returns a partial update. This makes the data flow explicit and, crucially, *checkpointable*: `LangGraph` can freeze the whole state to disk between turns (Section 9).

3.1 Reducers: how updates are merged

When a node returns Δs , each key k in Δs is merged into the state by a *reducer* ρ_k :

$$s'[k] = \rho_k(s[k], \Delta s[k]).$$

By default ρ_k is **last-write-wins** ($\rho_k(a, b) = b$). A field may instead declare a custom reducer via a Python Annotated type. Only two fields here use a non-trivial reducer:

- `search_results` uses `bounded_add`, an *append-with-cap*:

$$\text{bounded_add}(L, R) = (L \parallel R)[-limit:], \quad \text{limit} = 100,$$

i.e. concatenate the old and new lists but keep only the last 100 elements, so a checkpoint cannot grow without bound. This is the reducer used by the parallel web-search workers, each of which appends its own result.

- Every other field is last-write-wins.

Why the loop counters are *not* append-reduced. The bounded agent loop (Section 6) keeps a counter `loop_turn` and a list `observations`. These are deliberately plain last-write-wins channels. Because each new turn starts by writing `loop_turn=0` and `observations=[]`, a last-write-wins field *resets* cleanly at the start of every turn — even when the checkpointer restored the previous turn’s value. An append-style reducer would instead *merge across turns* (appending `[]` is a no-op), so the loop would never reset. The single node that writes them (gnodereflect) is sequential and returns the whole updated list each time, so last-write-wins is safe.

3.2 Field-by-field specification

Table 1 lists every channel of `AgentState`, its reducer, and its lifetime. “Per-turn” fields are reset each message by `create_initial_state`; “carried” fields are seeded from the conversation’s persisted state.

Table 1: The `AgentState` channels.

Field	Reducer / life-time	Meaning
<code>user_query</code>	last-write / per-turn	The composed prompt: long-term memory block + recent history + current message (Section 9).
<code>user_id</code>	last-write / carried	Stable identifier of the human. Namespaces long-term memory and durable preferences.
<code>session_id</code>	last-write / carried	The <i>conversation</i> id (one chat thread). Used as the checkpointer thread key.
<code>extracted_context</code>	last-write / per-turn	Side-channel facts for this message: the pristine <code>current_message</code> , the resolved <code>reply_language</code> , any “Ask-AI” focused-listing fields, previous results, comparison data.
<code>user_preferences</code>	last-write / carried	Coarse preference buckets: <code>hard_preferences</code> , <code>soft_preferences</code> , <code>excluded_areas</code> , <code>required_amenities</code> , <code>safety_concerns</code> .

Field	Reducer / life-time	Meaning
accumulated_search_criteria	last-write / carried	The structured search profile built up over the conversation: <code>area/areas</code> , <code>commute_destination</code> , <code>max_budget</code> , <code>max_travel_time</code> , <code>room_type</code> , <code>bedrooms</code> , <code>no_commute</code> , <code>move_in_date</code> , <code>criteria_gate_shown</code> , ...
tool_decision	last-write / per-turn	The chosen route: <code>{tool, params, reason, ...}</code> .
tool_observation	last-write / per-turn	The latest tool's output rendered as text (fed to the generator).
tool_raw_data	last-write / per-turn	The latest tool's structured payload (fed to the formatter and critic).
search_results	bounded_add (cap 100) / per-turn	Accumulator for the parallel web-search workers.
final_response	last-write / per-turn	The user-facing text.
response_type	last-write / per-turn	<code>answer</code> <code>question</code> <code>clarification</code> .
tool_data	last-write / per-turn	Structured payload for the frontend (listing cards, safety panel, POI list, ...).
run_id	last-write / per-turn	Fresh UUID per turn; seeds tool idempotency keys and tags the fan-out workers.
request_id	last-write / per-turn	Correlates logs/telemetry for one HTTP request.
context_tainted	last-write / per-turn	Taint bit: true once any <i>untrusted external text</i> (scraped listings, web results) has entered the state. Gates write-tools and triggers sanitisation (Section 12).
critic_attempts	last-write / per-turn	How many corrective regenerations the grounding critic has spent.
verdict	last-write / per-turn	The critic's structured judgement.
memory_context	last-write / per-turn	Reserved slot for a rendered memory block.
plan	last-write / per-turn	Reserved slot for a planned search list.
loop_turn	last-write / per-turn (resets to 0)	Count of loopable-tool executions completed this turn.
observations	last-write / per-turn (resets to [])	One record per loop step: <code>{turn, tool, observation, params_digest}</code> .

The three identity keys — keep them straight

user_id identifies a *person* and scopes durable memory and preferences. **session_id** identifies a *single conversation thread* (a chat) and is the LangGraph checkpoint key. A third key, **conversation_id**, is what the web layer calls the chat; here it is passed *into* the graph as **session_id**. The checkpoint thread is thus `f"{user_id}:{conversation_id}"` — so two users, or two chats of the same user, can never read each other's state. This precise scoping

is what fixed an early “memory bleed” defect.

4 The graph: nodes, edges, and the shape of a turn

The agent is the compiled LangGraph in `app/core/langgraph_agent.py`, function `build_agent_graph`. It has nine core nodes (plus three optional nodes woven in only when advanced features are enabled). Figure 1 shows the topology; the subsections then describe each node.




Figure 1: The compiled agent graph. Solid nodes are deterministic Python; only `DECIDE_TOOL` (intent classification), `REFLECT` (loop control), and `GENERATE_RESPONSE` (answer writing) call the LLM. Branch labels name the routing condition. The dashed advanced nodes `HYDRATE_PREFS`, `PERSIST_PREFS`, and `CONFIRM_SEARCH` are omitted for clarity.

4.1 The nine core nodes

extract_preferences. Cheap, deterministic. Scans the raw message for coarse preference cues (safety words, amenity keywords, exclusions, lifestyle signals) and folds them into `user_preferences`. It also applies any *explicit* in-message change to the accumulated search criteria — a new budget (“my budget is now £1000”), a city switch (“actually Manchester”),

a room type, a commute-time limit, a “no commute” declaration, or a move-in date — via `_apply_explicit_criteria_updates`, reusing the search tool’s regular-expression extractors so a standalone correction survives even in a turn that runs no search.

decide_tool. The router (Section 5). It computes a *decision* {`tool`, `params`, `reason`, ...} through a deterministic ladder ending in an LLM classifier, then jumps to the appropriate successor with `Command(goto=...)`: `direct_answer` → `GENERATE_RESPONSE`, `clarification` → `FORMAT_OUTPUT`, `multi_search` → `DISPATCH_SEARCHES`, any real tool → `EXECUTE_TOOL`.

execute_tool. Runs the chosen tool. It injects the accumulated criteria into the tool’s parameters (area, budget, commute destination, room type, bedrooms, move-in date, ...), attaches an idempotency key (Section 7), awaits `tool_registry.execute_tool`, and records the result as `tool_observation` (text) and `tool_raw_data` (structured). It then routes: a *loopable* tool that did not error → `REFLECT`; a structured card (safety, POIs, commute cost, listings) → `FORMAT_OUTPUT`; otherwise → `GENERATE_RESPONSE`. It also sets the `context_tainted` bit whenever the tool returned untrusted external text.

reflect. The bounded-loop controller (Section 6). After a loopable tool it decides “answer now” vs. “one more different tool”, subject to a hard turn cap and a no-progress guard, all failing *closed* to answering.

dispatch_searches / search_worker / gather_searches. The map–reduce subgraph for multi-query web research: fan out one worker per planned sub-search (LangGraph `Send`), run them concurrently, then fold their results into one combined observation.

generate_response. Calls the generation LLM to write the answer from the assembled prompt (security directive + language directive + synthesis template + observation). This is the only node that produces free prose.

critic. The grounding critic (Section 11). It checks the drafted answer’s monetary figures against the evidence and, if unsupported, triggers exactly one corrective regeneration; a persistently-failing answer is delivered with an appended caveat rather than being discarded.

format_output. The single choke point before `END`. It turns structured tool data into the frontend payload (`tool_data`: listing cards, safety panel, POI list, ...) and applies the final output sanitiser (Section 12) to *every* user-facing string.

4.2 Static vs. dynamic edges

Most edges are static (`graph.add_edge`): `EXTRACT_PREFERENCES` → `DECIDE_TOOL`; `GENERATE_RESPONSE` → `CRITIC` → `FORMAT_OUTPUT` → `END`; and the map–reduce chain. The three control-heavy nodes — `DECIDE_TOOL`, `EXECUTE_TOOL`, `REFLECT` — instead choose their own successor at runtime by returning `Command(goto=...)`. This is what lets one message flow through a variable number of tool calls while the graph’s *topology* stays fixed and analysable.

4.3 Three optional nodes

When advanced runtime flags are on, the builder weaves in `HYDRATE_PREFS`/`PERSIST_PREFS` (durable cross-conversation preferences, Section 9) and `CONFIRM_SEARCH` (a human-in-the-loop pause before the expensive fan-out). By default all three are disabled, so the production topology is exactly Figure 1.

5 Routing: how a message reaches the right capability

Routing is the intellectual centre of the system. The function `_compute_decision` inside `DECIDE_TOOL` is a *priority ladder*: a sequence of deterministic checks, each of which may short-circuit with a final decision, ending — only if nothing fired — in a single LLM classification call. The ordering is deliberate and load-bearing; each rung exists to stop a specific failure mode observed in testing.

Why deterministic-first?

An LLM classifier is excellent at open-ended intent but will occasionally re-run a fresh property search when the user asked “is the second one safe?”, or normalise a discriminatory request into a bland clarification, or answer a “don’t search yet, just research prices” message with a search form. Each such case is caught by a cheap, exact rule placed *above* the classifier, so the classifier only ever sees the genuinely ambiguous remainder.

5.1 The priority ladder

The rungs, in order (with the guard that fires each):

1. **Fair-housing screen (Equality Act 2010).** *Must be first.* If the *raw current message* asks to filter or avoid housing — or to gauge an area’s “safety” — by a protected characteristic (race, ethnicity, nationality, immigration status, religion, disability, family status, sexual orientation), the turn is refused deterministically with a bilingual message. Detection (`_fair_housing_violation`) requires an *avoidance operator* near a *protected-group term*, or a racial-purity selector — so positive preferences (“near a Chinese community”, “mosque nearby”) and crime-only “is E1 safe?” pass through untouched. Placing this above search routing means a search-shaped discriminatory request can never be laundered into a clarification.
2. **Memory-recall questions.** “What do you remember about me?” (and its Chinese equivalents) → `direct_answer`, answered from the injected long-term-memory block already at the top of `user_query`.
3. **Focused-listing context (“Ask AI”).** If the user clicked “Ask AI” on a specific card (`extracted_context['property_address']` is set) and the message is *not* a location/amenity/safety/commute question, answer from that listing’s own record (`reasoning_property`). Location-type questions escape to their proper tool.
4. **Follow-ups over the existing results.** With previous results in context: a comparative/-superlative question (“which of these is cheapest?”) → `direct_answer` over a comparison table; a detail question (“tell me more about the second one”) → `reasoning_property` over that one record; a suitability/opinion question (“would you recommend this for a couple?”) → the *listing-advice* route. These are answered from the conversation’s own real results, never a fresh search — fixing a class of bugs where such follow-ups bounced into “found 15 properties”.
5. **market-info negative guard.** An explicit “don’t search (yet), just research prices” instruction, or a research verb paired with a price-level noun and no listings request, is *market research* → the web-search synthesis path, never the search tool or its gate.
6. **Greetings** → `direct_answer`.
7. **Explicit “I don’t commute”** → search with the commute filter removed (never the endless “where do you commute to?” loop).

8. **Live-transport questions** (tube/train fares, journeys, travelcards, line status) → the dedicated TfL tool `get_transport_info`.
9. **Soft-gate follow-up.** If the recommended-criteria gate has already been shown this conversation and the user now supplies a budget/room-type/commute answer or says “continue”, proceed straight to the (confirmed) search.
10. **The LLM classifier** (`_majority_vote`). Only reached if nothing above fired.

5.2 The LLM classifier

The classifier is a *single* $T = 0$ call over a **static intent catalog** (`_INTENT_CATALOG`) of thirteen routes, each with a one-line description and three few-shot cues (English and Chinese). Crucially it classifies the *stripped current message* — the injected memory and history are passed only as secondary context “to resolve pronouns, never as the thing to classify” — because feeding the whole memory-prefixed query used to mis-route price-research questions into a property search. The model must return strict JSON `{"intent": "..."}` , parsed by a robust ladder:

strict JSON → embedded JSON in prose → longest-first substring match → heuristic fallback.

Two post-processing rules follow: a deterministic tie-break forces `search_properties` when unmistakable action verbs are present (“find me”, “帮我找房”); and if parsing fails entirely, the system prefers answering *over the shown listings* rather than defaulting to a fresh search (the root cause of the “found 15 properties” regression).

The static catalog and MCP

The intent catalog is hard-coded rather than derived from the tools’ own descriptions. This is deliberate: when tools are served over MCP (Section 7), an external client sees only names and schemas, not the project’s rich routing guidance. A static catalog makes classification quality *identical* whether tools come from the in-process registry or the MCP provider.

5.3 Reference resolution

Several routes need to know *which* listing the user means. `_resolve_last_result` and `_resolve_target_address` implement a uniform resolver: an explicit UK postcode in the message wins; then an ordinal (“the second one”, and the Chinese 第二个/第 2 套 — resolved by a dedicated regex because CJK has no word boundaries); then “#2 / number 2 / the last one”; then a deictic (“this one”, 这个房源) mapping to the first result; then a fuzzy match on listing name or street; finally the frontend-provided focused address. All resolvers read the *pristine* current message, never the memory/history-prefixed query, so a postcode from three turns ago cannot hijack this turn.

6 The bounded agent loop and web-research fan-out

Most turns call exactly one tool. But some questions genuinely need several — “is the area safe *and* what’s the commute cost *and* what’s nearby?”. For these the agent runs a **bounded loop**: after a *loopable* tool returns, a controller node decides whether the evidence already answers the question or whether one more, *different* tool call is warranted.

6.1 Loopable tools and the reflect controller

Only a fixed set of read-only tools may participate:

$$\mathcal{L} = \{\text{web_search}, \text{check_safety}, \text{search_nearby_pois}, \text{calculate_commute}, \text{calculate_commute_cost}, \\ \text{get_transport_info}, \text{get_weather}, \text{get_property_details}, \text{recall_memory}\}.$$

Everything else — a property search, the fan-out, a direct answer, or *any* tool that errored — keeps single-pass behaviour with zero added latency. After a loopable tool, EXECUTE_TOOL routes to REFLECT, which:

1. appends the just-run tool’s output to the per-turn **observations** list, tagged with a *parameter digest* (a stable hash of the tool name plus its non-volatile parameters);
2. asks a cheap $T = 0$ LLM whether to **answer** or **continue** with a named next tool and a refined sub-question;
3. if **continue**, builds that tool’s decision and loops back to EXECUTE_TOOL; otherwise routes to **answer**.

6.2 Three hard stops — all failing closed

The loop cannot run away. It terminates (routes to answering) whenever any of these fires:

- **Turn cap.** At most `MAX_AGENT_TURNS = 10` loopable executions per user message. The LangGraph `recursion_limit` is set to 80 — each loop iteration is two super-steps, so 10 iterations plus the fixed start/end nodes stays well under budget.
- **No-progress guard.** If the proposed next call has the *same* (tool, parameter-digest) as a call already made this turn, it is rejected — the loop can never repeat itself.
- **Any controller failure** — an unparseable reply, an exception, a non-loopable proposed tool, or a proposed call that degraded to a clarification — ends the turn by answering.

When more than one tool ran, the final answer is synthesised over *all* observations (each with a per-tool header), not just the last, and the grounding critic validates against every one of them.

Why “fail closed”? Every ambiguous situation resolves toward *stopping and answering with what we have*, never toward more tool calls. This guarantees termination and bounds cost, at the acceptable price of occasionally answering with slightly less evidence than a perfect agent might have gathered.

6.3 The map–reduce fan-out for web research

A market-research or general-knowledge question is planned into up to ten sub-searches by a planning LLM (`_plan_web_searches`), then executed as a LangGraph **map–reduce**:

- **Map** (`fan_out_searches`): emit one `Send` per sub-search, spawning `SEARCH_WORKER` instances that run *concurrently*. Each worker appends its result to the `search_results` channel via the `bounded_add` reducer (Section 3), tagged with the turn’s `run_id` so results from different turns never mix.
- **Reduce** (`GATHER_SEARCHES`): sort the workers’ results by index, fold them into one combined observation, and mark the context tainted (web text is untrusted). The combined evidence then flows to `GENERATE_RESPONSE`.

7 Tools: the agent’s hands

A tool is the only way the agent touches the real world. All tools share one contract, are registered in one place, and are made safe to retry through an idempotency ledger.

7.1 The uniform contract

Every tool is a `Tool` object (`app/core/tool_system.py`) carrying a name, a natural-language description (shown to the LLM), a JSON-schema parameter specification, an implementation `func` (sync or async), and metadata: `version`, `side_effect` $\in \{\text{none}, \text{write}\}$, and `retry_safe` (defaulting to “true iff `side_effect==none`”). Every call returns the same envelope:

```
@dataclass
class ToolResult:
    success: bool          # logical success
    data: Any = None       # structured payload
    error: str | None = None
    execution_time_ms: float | None = None
    tool_name: str | None = None
    version: str = "1"
    idempotency_key: str | None = None
```

Note the distinction between *logical* and *transport* success: a tool that returns `{'success': False, 'error': ...}` (the convention for a domain-level outcome such as “I need more information”) is surfaced as `ToolResult(success=False)` even though no exception was raised. This lets routing treat “need clarification” as data, not as a crash.

The `extra='ignore'` pitfall

A tool’s parameters are validated by a Pydantic model generated from its schema. Any argument the agent passes that is *not declared* in the schema is *silently dropped* before the implementation runs (Pydantic’s default `extra='ignore'`). This once broke the mid-conversation budget override: the agent injected `current_message`, but because it was undeclared it vanished. The rule, stated in-code, is now: *every parameter the router injects must be declared in the tool’s schema*.

7.2 The registry

`create_tool_registry()` builds the single canonical `ToolRegistry` with twelve tools. The registry owns per-tool statistics, a shared idempotency store, and the dispatch method `execute_tool(name, **kwargs)`. Async implementations are awaited; sync ones are pushed to a thread pool so they never block the event loop. Table 2 is the full catalogue with data sources.

Table 2: The twelve registered tools and their real data sources.

Tool	Purpose	Live data source
<code>search_properties</code>	Find real rental listings (the core tool).	OnTheMarket scraping (on-demand, sqlite-cached) + TfL commute annotation + FAISS/RAG re-rank.
<code>calculate_commute</code>	Travel time and route legs between two addresses.	TfL Journey Planner ; straight-line estimate outside London.

Tool	Purpose	Live data source
calculate_commute_cost	Time + fare + monthly cost.	TfL time; UK-postcode → zone inference; fares from a bundled TfL CSV.
check_safety	Crime/safety score for an address or area.	UK Police API (police.uk), last 6 months.
check_transport_cost	Official 2025 Travelcard/PAYG prices by zone.	Hard-coded TFL_FARES_2025 table.
get_transport_info	Live journeys, fares, travelcards, line status.	TfL Unified API ; London-only, honest “outside London” outside the bounding box.
get_weather	Current weather for a city.	Open-Meteo (no key required).
search_nearby_pois	Nearby amenities (supermarket, gym, station, ...).	OpenStreetMap Overpass + Nominatim geocoding.
get_property_details	Facts/policies about one already-shown listing.	Local property CSV database.
web_search	General web questions; also a sub-query coordinator.	Web snippet provider (DuckDuckGo/SearXNG-style).
recall_memory	Read the user’s long-term memory.	Local ChromaDB (Section 8).
remember	Write a durable fact (<i>the only write tool</i>).	Local ChromaDB; side_effect="write".

7.3 Idempotency: making retries safe

A tool call that times out or is replayed after a crash must not perform its side effect twice. The mechanism is a content-addressed key plus a durable ledger.

The key is deterministic (`ToolInvocation.create`):

```
canonical = json.dumps(params, sort_keys=True, ensure_ascii=False, default=str)
idempotency_key = sha256(f"{run_id}:{node_id}:{tool}:{canonical}").hexdigest()
```

Because `sort_keys=True`, the key is invariant to parameter ordering; the same (run_id,node_id,tool,params) always hashes identically. The ledger (`IdempotencyStore`, a SQLite table) supports `claim` (atomic “I am running this”), `complete` (store the result), and `get` (return a prior completed result). For the one write tool, `execute` requires a key, returns a cached result on a repeat, and only runs the effect if it wins the `claim` race. The result is *at-most-once on effect, exactly-once on observed result*.

7.4 search_properties in depth

The core tool is a large async function whose schema declares *no required parameters* — it degrades gracefully, embodying the “zero workflow prior” principle: a user need not answer a fixed questionnaire. Its behaviour:

- **Memory stripping.** Before extracting hard filters it removes the injected `What I remember about this user:` block, so a filter is never pulled from a previous conversation’s remembered criteria.
- **This-turn overrides.** A budget, commute limit, room type, move-in date, or “no commute” stated *this turn* overrides the accumulated value, via the deterministic extractors below.

- **The hard area gate** (the *only* blocking gate). If no living area can be resolved — even after regex extraction, a natural-language fallback, and destination-derived defaults — the tool returns `status="need_clarification"` with `clarification_kind="missing_area"`.
- **The soft criteria gate**. The *recommended* fields are budget, room type, and commute. If the area is known but some recommended field is missing, the gate fires *at most once per conversation* (`criteria_gate_shown` is sticky-true), asking a single friendly question with a “search anyway” bypass. A move-in date is purely optional — it never triggers the gate alone; it only “rides along” when the gate fires for a recommended field.
- **Soft budget filtering**. Within budget is a perfect match; up to $1.15\times$ budget is a flagged soft violation; above is excluded. A weekly budget is normalised to monthly by $\times 4.33$.

Deterministic bilingual extractors

A family of regular-expression extractors is the shared source of truth for parsing, reused by the router, the preference node, and the tool itself. Each handles English *and* Chinese:

- `_extract_budget` $\rightarrow$ (amount, period): £1800, 1800 pcm/pw, “budget of 1800”, “under/around 1200”, and Chinese “1000 以内/左右/预算”; sanity range 200–20000.
- `_extract_area`: Chinese city names via a ~ 40 -entry map (曼彻斯特 $\rightarrow$ Manchester), or a switch cue (“make it ...”, “flats in ...”) validated against a curated UK-place list; deliberately conservative (“Bath”/“Reading” as bare words never trigger a switch; Mars/火星 $\rightarrow$ none).
- `_extract_room_type` $\rightarrow$ studio | ensuite | shared, via a bilingual synonym table.
- `_extract_commute_minutes`: “half an hour” $\rightarrow$ 30, “an hour” $\rightarrow$ 60, “N minutes”, “N hours”; clamped 1–180.
- `_extract_no_commute`: “wfh”/“work from home”/“just living” and Chinese 不通勤/在家办公.
- `_is_proceed_intent`: “continue”/“search anyway” and 继续/直接搜.
- `_extract_move_in_date` $\rightarrow$ YYYY-MM-DD: full dates always fire; a bare month resolves to the first day of its next occurrence; guarded so the modal “may” never becomes the month May.

7.5 The model router

Different graph nodes need different model behaviour, so a small `ModelRouter` (`src/uk_rent_agent/llm/router.py`) maps a *purpose* to a concrete DeepSeek model, temperature, and token budget:

Purpose	Model	Temp.	Used by
intent	deepseek-chat	0.0	the router’s classifier
planner/critic	chat, or reasoner if complex	0.0	search planning, loop controller
responder	deepseek-reasoner	0.1	GENERATE_RESPONSE
memory/judge	deepseek-chat	0.0	memory extraction, grading

Classification and grounded generation run near-deterministic ($T \approx 0$); only the reasoning responder is given room to think. The provider is swappable to a local Ollama server via an environment variable, so the whole system can run offline.

8 Memory: what the agent remembers, and for how long

Memory is the feature users most anthropomorphise (“it remembers my budget!”), and the one with the sharpest correctness and privacy requirements. The system has two architecturally separate memory families plus the graph’s own short-term state.

8.1 The three time-scales

Short-term (within a turn). The `AgentState` (Section 3). It exists only while one message is being processed and is not a persisted “memory” at all — it is the live working set.

Medium-term (within a conversation). The per-conversation *checkpoint* and the in-memory session cache. This is what makes “the second one” resolvable three turns later. Keyed by the conversation thread (Section 9).

Long-term (across all of a user’s conversations). The layered `AgentMemory` store (below) and the optional durable preference profile. Keyed by `user_id`, so a fact learned in one chat surfaces in another.

8.2 The layered long-term memory

`AgentMemory` (`app/rag/agent_memory.py`) is a single ChromaDB collection, `agent_memory`, holding three kinds of item distinguished by a metadata field `mtime`. The design follows the *Generative Agents* and *Mem0* lines of work, with this codebase’s own constants:

Episodic — individual events (“User asked: ... [assistant used `check_safety`]”). Written every turn.

Semantic — distilled durable facts about the user (budget, destination, constraints). Extracted from each turn and *consolidated*.

Reflection — higher-level insights synthesised periodically from the accumulated memories.

Writing: extraction and consolidation (Mem0-style)

After each turn, on a background thread, `remember_turn`: (1) writes an episodic record; (2) calls an LLM to extract up to eight atomic durable facts as strict JSON; (3) *consolidates* each fact against the four nearest existing semantic memories, letting the LLM emit `ADD` / `UPDATE` / `DELETE` / `NOOP` operations — so a refined budget *replaces* the old one rather than piling up, and a contradiction deletes the stale fact. An MD5 hash of (`mtime` | `user_id` | `text`) guards against exact duplicates.

Importance and the reflection loop

Each memory carries an *importance* in `[1, 10]`: for episodic memories it is an LLM “poignancy” rating (a hard budget or fixed destination scores high; small talk low), defaulting to 5; semantic and reflection memories take fixed defaults 7 and 8. Importance accrues per user; when the accumulated importance crosses a threshold of 30, a *reflection* pass distils one-to-three insights (importance 8) from the user’s last thirty memories and resets the counter. This is the Generative-Agents reflection mechanism, scaled down from the paper’s game-day threshold of 150.

Retrieval: the Generative-Agents score

Given a query, retrieval fetches the top $K = 25$ candidates by vector similarity (filtered to this `user_id`), then re-ranks each candidate m by three signals:

$$\begin{aligned}\text{relevance}(m) &= 1 - \text{cosine_distance}(q, m), \\ \text{recency}(m) &= 0.995^{\Delta t_{\text{hours}}(m)} \quad (\Delta t \text{ from last access}), \\ \text{importance}(m) &= \text{rating}(m)/10.\end{aligned}$$

Each signal is independently min-max normalised to $[0, 1]$ across the candidate set, and the final score is their *unweighted sum*:

$$\text{score}(m) = \widehat{\text{relevance}}(m) + \widehat{\text{recency}}(m) + \widehat{\text{importance}}(m) \in [0, 3].$$

The top $n = 6$ are returned and their `last_access` is bumped to now — so retrieving a memory makes it more retrievable next time (the paper’s reinforcement effect). The embedding model is `all-MiniLM-L6-v2` (384-dimensional); property retrieval additionally uses an explicit FAISS inner-product index over L2-normalised vectors (equivalent to cosine).

Privacy: fail closed by design

Every read and write is gated by `_valid_user_id`, which rejects a blank id *and the literal “default”* — the implicit shared namespace that once let one user’s budget leak to the next. With no valid id, writes are refused and reads return empty; there is never a global fallback. The Chroma query always carries `where={"user_id": id}`; legacy records written before namespacing carry no id and so can never match a user-scoped filter (inert, not leaked). `forget(user_id)` is the GDPR erasure path, wired to `/api/forget_me`.

8.3 Memory as tools

The agent can also read and write memory explicitly, as tools: `recall_memory` (returns both structured memories and a ready-to-paste prompt block) and `remember` (writes a fact of a given kind). Both re-check `_valid_user_id` at the tool boundary — defence in depth on top of the store’s own gate — and `remember` is the system’s only *write* tool, hence the only one subject to the idempotency ledger and the taint gate.

9 From HTTP request to agent turn: the web and session layer

The agent is served by a Flask application (`app/app.py`). This section traces one request end-to-end and explains how many users and many conversations stay isolated.

9.1 The request lifecycle

The main conversational endpoint is `POST /api/alex`. In order:

1. **Middleware.** Optional authentication enforcement, then per-path *rate limiting* (a process-local sliding-window limiter: 12/min for `/api/alex`, 20/min for `/api/search_direct`, ...).
2. **Parse and validate** the JSON body; a non-string `conversation_id` is rejected *before* it can reach SQLite.
3. **Identity resolution** (`resolve_identity`). A priority ladder: an authenticated session wins; otherwise, only if explicitly allowed, a client-supplied id (validated against `^[A-Za-z0-9_-]{1,64}$`); otherwise the signed session cookie; otherwise a freshly minted UUID. A logged-in session is *authoritative* — a client cannot impersonate an account by spoofing a header.

4. **Conversation lookup** or implicit creation (title derived from the message, HTML-stripped as a stored-XSS defence). The user turn is persisted *up front* so it survives a crash mid-generation.
5. **Compose the query** (below).
6. **Build AgentState** and invoke the graph asynchronously with a config whose thread key is `f"{user_id}:{conversation_id}"` and whose recursion limit is 80.
7. **Extract** `final_response` / `response_type` / `tool_data` and **persist**: write back the conversation's cached state, durably store the assistant message, and fire a background long-term memory write.

Agent-side outcomes always return HTTP 200 (an “error” response type is a normal renderable payload), so the frontend can adopt the freshly minted `conversation_id`.

9.2 Composing the query: layered context injection

The single string handed to the graph as `user_query` is built in layers, *prepended* so the model reads them top-down:

```
What I remember about this user:      <-- long-term memory, keyed by user_id
[FACT] budget is about GBP 1500 pcm
[INSIGHT] prioritises safety near campus

Previous conversation:                <-- last 3 turns (or 5 for a clarification answer)
User: ...
Alex: ...

Current user message: <the raw message>
```

Two invariants make this safe. First, the *pristine* message (no memory/history prefix) is stashed separately at `extracted_context['current_message']`, so tools that parse “only this message” (budget/postcode/ordinal regexes) never scrape a stale value out of injected memory. Second, long-term memory is retrieved and namespaced by `user_id` — shared across all of that user’s conversations, but never across users.

9.3 Two isolation axes, two persistence mechanisms

There are effectively two identity axes (a naming caveat: the code’s `session_id` is a vestigial alias — the graph receives the `conversation_id` under that parameter name):

Axis	Identifies	Scopes
<code>user_id</code>	a person, across logins/de-vices/chats	long-term memory, durable preferences, favourites, ownership of all conversations
<code>conversation_id</code> (= <code>thread_id</code> part)	one chat thread	the transcript, the LangGraph checkpoint, the session hot-cache slice

The two persistence mechanisms are distinct:

- **Short-term, per-conversation** — a LangGraph *checkpointer* (a thread-safe SQLite saver) keyed by `f"{user_id}:{conversation_id}"`. It holds the full per-conversation graph state and enables pause/resume and time-travel.

- **Long-term, per-user** — an optional cross-thread `BaseStore` keyed by the namespace `("user_prefs", user_id)`. Because the namespace is per-user, a preference saved in conversation A is visible in conversation B — the whole point of the store-vs-checkpointer split. Two graph nodes, `HYDRATE_PREFS` (entry) and `PERSIST_PREFS` (exit), read and write it, persisting only the *profile* (structured criteria), never the transcript.

How the “memory bleed” bug was fixed

The in-conversation working state was once bare module globals shared by every caller, so one chat’s budget could surface in another. It is now a `SessionStore` keyed by the tuple `(user_id, conversation_id)`, each slice a fresh deep copy, each guarded by its own per-key turn lock. Durable preferences additionally *fail closed* on sentinel ids `{"default", "anonymous"}`, so two people who both present “default” can never share a bucket. Every conversation route checks ownership and returns 404 on a mismatch.

9.4 Concurrency: snapshot–compute–commit

A single conversation may receive two overlapping requests (two browser tabs). Each turn (i) snapshots the cached state under the per-conversation lock, (ii) releases the lock across the slow LLM call, then (iii) re-acquires it to append the result *in place*. So a concurrent turn’s append is never clobbered, yet the expensive generation never holds the lock.

10 Ranking listings: a multi-objective utility

Once candidate listings are retrieved, they must be ordered. The ranker (`src/uk_rent_agent/agent/ranking.py`) is a transparent, hand-tuned utility — deliberately not a learned black box, so its behaviour is auditable.

10.1 Hard constraints, then soft scoring

First, *hard filters* drop (not penalise) any listing that violates a non-negotiable constraint:

- over budget: `price > max_budget` $\Rightarrow$ dropped with reason `over_budget`;
- in an excluded area: the address contains an excluded-area substring $\Rightarrow$ dropped with reason `excluded_area`.

Each surviving listing ℓ then receives a score

$$\text{score}(\ell) = \underbrace{1.0}_{\text{base}} + \underbrace{\max\left(0, 1 - \frac{\text{price}(\ell)}{B}\right)}_{\text{budget affinity}} + \underbrace{0.05 \cdot \mathbf{1}[\text{source} = \text{local_faiss}]}_{\text{provenance}},$$

where B is the budget. The budget-affinity term is a normalised “cheaper is better” objective in $[0, 1]$: a listing at price 0 contributes +1, one at exactly budget contributes +0 (and, since over-budget listings are already dropped, the term never goes negative). The provenance term gives a small edge to the trusted local index. Listings are sorted by score descending and truncated to a limit (default 10). For a budget-known listing the score lies in $[1.0, 2.05]$.

This is a *linear additive utility* $U = w_0 + w_1x_1 + w_2x_2$ with a constant base and two soft objectives, filtered by two hard constraints. The weights $(w_0, w_1, w_2) = (1, 1, 0.05)$ are fixed and legible; there is no training, so a product owner can reason about exactly why one listing outranks another. (A separate RAG re-ranker used inside the search tool adds semantic-similarity, commute, and safety-trend terms with weights 0.4/0.3/0.2/0.1 for the

in-tool ordering.)

11 The grounding critic: catching hallucinated numbers

An LLM asked to summarise listings will occasionally state a rent that appears nowhere in the data. The critic (`src/uk_rent_agent/agent/critic.py`) is a deterministic guard that catches fabricated *monetary figures* while never destroying a legitimate answer.

11.1 What “grounded” means, precisely

The critic does not compare strings; it does arithmetic. It extracts every *money mention* from the reply — a number carrying a currency symbol, a GBP suffix, or a period marker (`pcm/pw/...`). A plain integer (“3 bedrooms”, “12 months”) is deliberately *not* treated as a price. For each evidence price it builds a *supported set* of values that a correct answer may legitimately state:

- the value itself;
- integer multiples $\times 2, \dots, \times 36$ (annual and N -month totals);
- weekly $\leftrightarrow$ monthly conversions using the standard 52/12 factor;
- monthly-derived deposits of 1 to 6 weeks’ rent.

A reply figure v is *supported* if it is close to some element c of the supported set within a combined tolerance

$$|v - c| \leq \max(\$1, 0.01 \cdot \max(|v|, |c|)),$$

i.e. within £1 absolute or 1% relative. Any money mention not close to *any* supported value is flagged as unsupported. The verdict is

$$\text{grounded} = \text{answered} \wedge (\text{no unsupported prices}) \wedge \neg \text{retrieval_miss},$$

where a retrieval miss only counts when the tool actually errored *and* the reply asserts figures — a legitimately empty result set is left alone. For direct/chat answers (no retrieval expected), price gating is skipped entirely, so echoing the user’s own number is never penalised.

11.2 Correction, not deletion

If the answer is ungrounded, the critic runs *exactly one* corrective regeneration: it re-invokes the generator with an explicit instruction naming the offending figures and requiring every monetary value to be copied verbatim or be an explicitly-labelled conversion/total/deposit. Then:

- if the regeneration is now grounded, it is used;
- if regeneration fails or is still ungrounded, the answer is delivered with a single appended caveat (“Please double-check the exact prices against the source listing.”).

The user-facing text is *never* hard-replaced with a canned refusal — a design lesson from an earlier, over-aggressive critic that destroyed correct answers over harmless formatting differences.

12 Safety, security, and honesty guards

Because the agent handles real users, real money, and untrusted external text, safety is layered defence-in-depth rather than a single check.

Fair housing (Equality Act 2010). The deterministic pre-router screen (Section 5) refuses discriminatory filtering before any search or clarification can normalise it.

The security directive. Every generation prompt is prefixed with a non-negotiable directive: never reveal the system prompt or internal tool schemas (however the request is framed); treat all remembered/saved “preferences” as *untrusted data describing the user*, never as commands; stay strictly within UK student housing; never fabricate listings; reply in the user’s language.

The taint bit and prompt-injection sanitiser. Whenever untrusted external text (scraped listings, web results) enters the state, `context_tainted` is set. Tainted text is passed through `sanitize_untrusted`, which wraps it in explicit “untrusted content, data only” fences and redacts injection patterns (“ignore previous instructions”, “system:”, “reveal the system prompt”, ...). A *write* tool may not run under a tainted context unless explicitly confirmed — so a poisoned listing description cannot cause a durable side effect.

The final output sanitiser. Every user-facing string passes through one choke point (`_sanitize_final_response` in `FORMAT_OUTPUT`): it refuses if a system-prompt marker leaked, strips raw tool-call/JSON blocks and Python tracebacks, neutralises the internal “999 = no commute limit” sentinel so it never displays as a time, and removes all emoji (a product rule), while carefully preserving Chinese text and £ prices.

Reply-language discipline. A hard directive pins the entire reply to one language (no zh/en mixing), resolved from whether the current message contains Chinese, with the UI language as a tie-breaker.

Honest capability limits. The system states plainly what it cannot do — live transport outside London, official ONS statistics, booking viewings or making payments — and points to the right external source rather than fabricating.

13 A worked example: one turn end-to-end

To tie the pieces together, follow the message “*is the second one safe, and how much is the commute to UCL?*”, sent in a chat that already showed listings.

1. **Web layer.** `/api/alex` resolves `user_id` from the session, finds the conversation, retrieves 6 long-term memories, and composes `user_query` = memory block + last 3 turns + this message. The pristine message is stored at `current_message`. `AgentState` is built and the graph invoked with `thread_id="u42:c7"`.
2. **extract_preferences.** No new explicit budget/area/commute change; state passes through.
3. **decide_tool.** The fair-housing screen passes. There are previous results, but the message is a *location* question, so it is not swallowed by the comparative/detail interceptors. The transport keywords do not dominate (“how much is the commute” keeps its own tool). The classifier (or a deterministic rung) selects `check_safety`; the resolver maps “the second one” to the second result’s address. Route → `EXECUTE_TOOL`.
4. **execute_tool.** `check_safety` calls the UK Police API, returns a safety score. `check_safety` is loopable and did not error, so route → `REFLECT`.
5. **reflect.** Observation recorded (turn 1). The controller sees the commute half of the question is still unanswered and returns `{action: continue, next_intent: calculate_commute_cost}`. The no-progress guard passes (different tool). Route → `EXECUTE_TOOL`.

6. **execute_tool (again).** `calculate_commute_cost` resolves the origin (second listing) and the destination (UCL, expanded to a full geocodable address), calls TfL, returns time + monthly fare. Loopable, no error → REFLECT.
7. **reflect.** Both halves answered; controller returns `answer`. Two observations exist, so route → GENERATE_RESPONSE.
8. **generate_response.** Builds the synthesis prompt (security + language directives + *both* observations with per-tool headers) and writes one coherent reply.
9. **critic.** Extracts the monthly fare figure, checks it against the commute-cost evidence within tolerance — grounded — passes unchanged.
10. **format_output.** The multi-tool synthesis is kept as prose (no single card can represent it); the final sanitiser strips any stray artifacts and emoji. END.
11. **Web layer.** The reply and updated criteria are written back under the turn lock; a background thread updates long-term memory (“asked about safety and commute of listing #2”).

In one message the agent made two real API calls, in the right order, kept the answer grounded, and never once risked an unbounded loop or a cross-user leak — which is precisely what the architecture was built to guarantee.

Glossary

Node A function in the graph that reads state and returns an update.

Edge A permitted transition between nodes; static or chosen at runtime.

Reducer The rule that merges a node’s update into a state field.

Route The decision of which capability handles a message.

Tool External-world code returning a uniform `ToolResult`.

Taint A flag that untrusted external text has entered the turn.

Idempotency key A content hash making a retried write safe.

Checkpointner Per-conversation persisted graph state (thread-keyed).

Store Per-user durable preference profile (namespace-keyed).

Grounding The property that every stated figure traces to evidence.

LangGraph The library expressing the agent as a graph over shared state.